

CSCI 5521: Machine Learning Fundamentals (Spring 2024)**Quiz 1 (Thursday, Feb 8)****Due on Gradescope at 02:00 PM, Friday, Feb 9****Instructions:**

- This quiz has 3 questions, 30 points, on 1 page.
 - Please write your name & ID on your submission pages.
- (6 points)** Supervised learning has a wide range of applications. For example, with the popularity of wearable devices, such as FitBit and Apple Watch, one of their core functionalities is to identify whether a time period contains physical activity or rest based on **heart rate data** collected from the device. Please model this problem as a two-class classification task and answer the following questions.
 - Name two relevant features to this two-class classification task (i.e., whether a time period contains physical activity or rest). Any reasonable features are acceptable.
 - What are the labels in this two-class classification task?
 - (10 points)** What could we do to reduce overfitting in a polynomial regression model? Select all the option(s) that apply.
 - Collect 30% more new training data.
 - Sample half of the original training data and **only** use that as the new training data.
 - Double the size of the original test data (by collecting new data) and replace as the new test data.
 - Increase polynomial degree.
 - Decrease polynomial degree so that the model becomes a linear regression model.
 - (14 points)** The conditional probability density functions of two classes C_1 and C_2 are shown in the figure below, with $P(x|C_1) \sim \mathcal{N}(1.5, 0.8)$ and $P(x|C_2) \sim \mathcal{N}(0.9, 1.3)$.
 - Assuming the priors are equal, predict which class (C_1 or C_2) the data point $x = 0.5$ (illustrated with the green dot) belongs to. Briefly explain why.
 - What if the priors are $P(C_1) = 0.6$ and $P(C_2) = 0.4$, respectively?

